



TF-IDF

Term Frequency / Inverse Document Frequency

LING 575

HW1

Einar Horn, Lu Liu, Avijit Vajpayee, Daniel Campos, Peter Schoener

Method introduction



TF-IDF

- A statistical method to represent how important an individual word is given a corpus of documents.
- The metric is a combination of components
 - **Term Frequency** : How common is the term in the given document
 - **Inverse Document Frequency** : How uncommon is the term across documents
- NOT a representation but a metric / numerical statistic
- Used as a weighting factor for representations of words in a corpus




Origin

- Inverse Document Frequency was first proposed by [Karen Sparck Jones](#) in 1972
 - First called term specificity
 - Proposed as a heuristic but proving theoretical foundations has been tricky
 - Loose Connection proposed to Zipf's law.
- Term Frequency was first proposed by [Hans Peter Luhn](#) in 1957



Zipf's Law

- Zipf's Law:
 - Many types of data studied in the physical and social sciences can be approximated with a Zipfian distribution, one of a family of related discrete power law probability distributions.
- Zipf's Law for NLP:
 - The frequency of any word is inversely proportional to its rank in the frequency table.
 - The frequency of the n -th most frequent word is roughly proportional to $1/n$
- Brown Corpus:
 - Thus "the" constitutes nearly 7%, "to" and "of" more than another 3% each; About half the total vocabulary of about 50 000 words are *hapax legomena*: words that occur only once in the corpus





Zipf's Law Takeaways

- The most frequent words are not informative
 - *a, and, the, be, of, with* etc.
- Long tail of rare words that occur just once
- Informativeness : balance between two ends



Relevancy to 575

- TF-IDF is not a text representation in the modern sense but represents of a words salience for a given document in respect to a corpus.
- Frequency used as a weighting factor
- This representation allows for a scalable way of understanding relative relations of words in a corpus in a scalable way.
- TF-IDF can scale to virtually any size corpus with minimal memory requirements.

Algorithm details



Basic Equation

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

t = term

d = document

D = corpus



Term Frequency

A count of how many times a term t appears in document d with various normalization functions

- Binary

$$0/1$$

- Raw Count

$$f_{t,d}$$

- Normalized

$$\frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$

- Log Scaled

$$\log(1 + f_{t,d})$$

- Adjusted for Document Length

$$0.5 + 0.5 \frac{f_{t,d}}{\max_{t' \in d} f_{t',d}}$$



Inverse Document Frequency

A measure of how much information the word provides (Is it common or rare across all documents).

n_t : # documents with term t

N : # documents

- IDF

$$\log \frac{N}{n_t}$$

- Smooth

$$\log \frac{N}{1+n_t}$$

- Max Normalized

$$\log \frac{\max_{\{t' \in d\}} n_{t'}}{1+n_t}$$

- Probabilistic

$$\log \frac{N-n_t}{1+n_t}$$



TF-IDF Example

- “the”
 - $tf(\text{“the”}, d_1) = 5/7$
 - $tf(\text{“the”}, d_2) = 6/9$
 - $idf(\text{“the”}, D) = \log(2/2) = 0$
 - $tf-idf(\text{“the”}, d_1) = tf-idf(\text{“the”}, d_2) = 0$
 - “The” is not an informative term
- “assignment”
 - $tf(\text{“assignment”}, d_1) = 0/7$
 - $tf(\text{“assignment”}, d_2) = 2/9$
 - $idf(\text{“assignment”}, D) = \log(2/1) = 1$
 - $tf-idf(\text{“assignment”}, d_1) = 0$
 - $tf-idf(\text{“assignment”}, d_2) = 2/9$
 - “Assignment” is more informative in Document 2

Document 1

Word	Freq
the	5
homework	2
assignment	0

Document 2

Word	Freq
the	6
homework	1
assignment	2


Performance



Count vectors vs TF-IDF vectors

- Task: Classify text documents by topic using various classifiers
- Compare results of using count vectors and tf-idf vectors
- Corpus: 18000 newsgroups posts on 20 topics
 - Topics: baseball, computers, religion, etc

Classification accuracy for count vectors vs tf-idf vectors



Takeaway: TF-IDF outperforms count vectors in every classifier

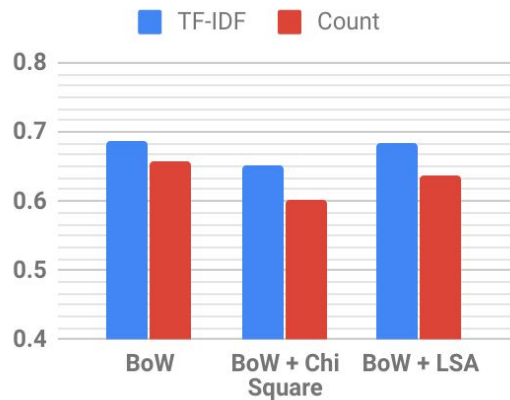


Space and Time Complexity

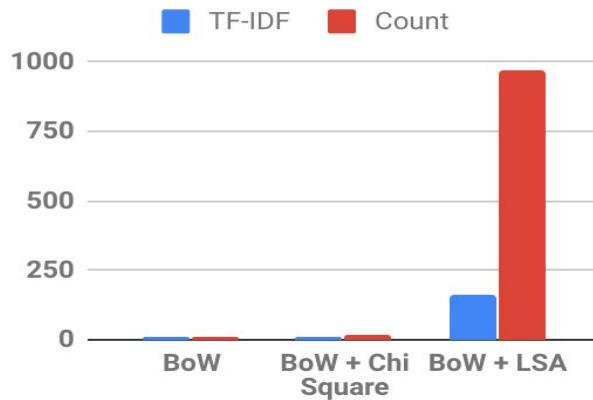
- **Count Vectors**: Single Pass over all documents , $O(V * D)$ space
- **TF IDF Vectors**: Two Passes over all documents , $O(V * D)$ space
- Collecting counts/collocation statistics : Fully parallelizable
- Space Complexity a bigger issue
 - Large sparse vectors not as ideal for “most” downstream ML estimators
 - Sparse vectors generally used with dimensionality reduction techniques
 - Top Terms selection through Chi-Square
 - Matrix Factorization (SVD / NMF)
 - TF-IDF weighted average of neural word embeddings

- TF-IDF vs. Count Sparse Vectors in 3 settings
- Linear SVM Classifier
- 20 News Groups Dataset

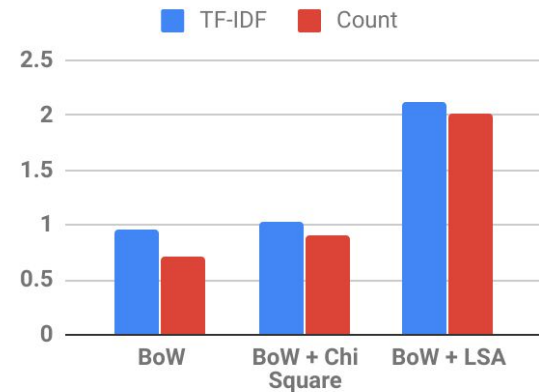
Accuracy



Training Time



Testing Time



Code : https://github.com/einarhorn/tf-idf-classification/blob/master/lsa_weighting_comparison.py

Applications



Stopwords Filtering

- TF-IDF to filter stopwords
 - Frequent words are weighted higher.
 - If the frequent words are frequent in all documents, then their weight is lowered.
 - Stopwords tend to have low IDF and TF-IDF.



Search Engines

- Ranking function, to rank matching documents according to their relevance to a given search query, simplest one is to sum the TF-IDF for each query term.
- The cores of ranking function - BM25 and TF-IDF.
- BM25 is the next generation of TF-IDF: it improves upon TF-IDF.
 - $f(q_i, D)$ is q_i 's term frequency in document D , $IDF(q_i)$ is the IDF weight of the query term q_i .
 - $|D|$ is the length of the document D in words.
 - Avgdl is the average document length in the text collection from which documents are drawn.
 - k_1 and b are free parameters, usually k_1 is between $[1.2, 2.0]$ and $b = 0.75$.

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)}$$



Recommendation Engines

- Content based recommendation system
- System recommends similar items according to keywords or properties of items with a distance measure.
 - Convert texts or words into vectors
 - TF-IDF to weight and construct a TF-IDF matrix.
 - Cosine similarity



Text Summarization

- Assign scores based on TF-IDF to sentences, and take the top scoring n sentences as a summary
- How to assign score to sentence?
 - Sum of TF-IDF values
 - Paper by Seki suggests the following:
 - sum noun terms
 - Biase the weights of words in sentence that are in title of the document
 - multiply this sum by the position value of the sentence in the document



Document Classification and Clustering

- Classify or cluster the documents - numerical representation of the sentences or documents.
- Represent the documents into mutually comparable vectors.
- The documents can be represented using the TF-IDF scores.
- Then we can represent each document as a vector of terms using a global ordering of each unique term found throughout all the documents.

Variants



Proportional Document Frequency (PDF)

- Measures the difference of how often a term occurs in different domains.
- Proportional Document Frequency is computed as $\exp(n_{jc}/N_c)$

$$W_j = \sum_{c=1}^{c=D} |F_{jc}| \exp\left(\frac{n_{jc}}{N_c}\right),$$

$$|F_{jc}| = \frac{F_{jc}}{\sqrt{\sum_{k=1}^{k=K} F_{kc}^2}},$$



IDuF

- The TF component in TF-IDuF is the same with TF-IDF.
- Classic IDF is calculated using the document frequencies in the recommendation corpus.
- IDuF is calculated using the document frequencies in a user's personal document collection c_u .

Sparse Word Vector Representations



Information Theoretic Background

Information Content / Self Information : (for an event) amount of information gained when sampled

$$I(E) = -\log(\Pr(E))$$

Entropy (H) : Expected value of information content for a random variable

Mutual Information : amount of information about one R.V. through observing another R.V.

$$MI(X, Y) = \sum p(x, y) \log \frac{p(x, y)}{p(x)p(y)}$$



Pointwise Mutual Information

- *Strength of association* : word to word ([Church and Hanks 1990](#))
- Distributional Hypothesis of Semantics
- Define a window of context length C
- Each word represented as a sparse vector of length |V| (size of vocab)
 - How many times each word in vocab appears within context of given word
 - Mutual Information between 2 R.V. is the expected value of PMI

$$PMI = \log_2 \frac{P(word_1, word_2)}{P(word_1) * P(word_2)}$$

$$PPMI = \max(PMI, 0)$$



Extensions to PMI

- Normalized PMI ([NPMI](#)) :

- Bound between $[-1, +1]$

$$\frac{PMI(w,c)}{I(w,c)} = \frac{PMI(w,c)}{-\log(p(w,c))}$$

- Context Smoothing ([Vector Semantics](#)) :

- PMI biased towards infrequent events
- Slightly shifts probability to rare context words

$$PPMI_{\alpha}(w,c) = \max(\log \frac{P(w,c)}{P(w)P_{\alpha}(c)}, 0)$$

$$P_{\alpha}(c) = \frac{count(c)^{\alpha}}{\sum_c count(c)^{\alpha}} \quad \alpha = \frac{3}{4}$$

- Shifted PPMI

- Related to SkipGram model of Word2Vec trained with Negative Sampling

$$SPPMI(w,c) = \max(PMI(w,c) - \log k, 0)$$



Connection to Dense Neural Embeddings

- [GloVe](#)
 - Objective directly minimizes the difference between dot product of word and context vectors and log of co-occurrences.
- [word2Vec : Neural Word Embeddings as Implicit Matrix Factorization \(Levy et. al. 2014\)](#)
 - SkipGram with Negative Sampling (SGNS) implicitly factorizes word context matrix
 - SGNS superior for word analogy
 - SVD on PMI matrix at least as good for word similarity
 - Excellent Blog - [Sebastian Ruder](#)


Appendix



Relevant Papers

- [Understanding inverse document frequency: on theoretical arguments for IDF](#)
- [A Statistical Approach to Mechanized Encoding and Searching of Literary Information*](#)
- [Sentence Extraction by tf/idf and Position Weighting from Newspaper Articles](#)
- [From RankNet to LambdaRank to LambdaMART: An Overview](#)
- [Relevance weighting of search terms](#)
- [A STATISTICAL INTERPRETATION OF TERM SPECIFICITY AND ITS APPLICATION IN RETRIEVAL](#)

Contributions

- 
- [Daniel Campos](#) : Introduction to TF-IDF, Zipf's Law, Variants
 - [Einar Aleksander Horn](#) : TF-IDF Calculations, Performance Benchmarking, Text Summarization
 - [Lu Liu](#) : Applications, Variants
 - [Avijit Vajpayee](#) : TF-IDF Formulations, Performance Benchmark, PPMI
 - [Peter Schoener](#) : Applications, PPMI