

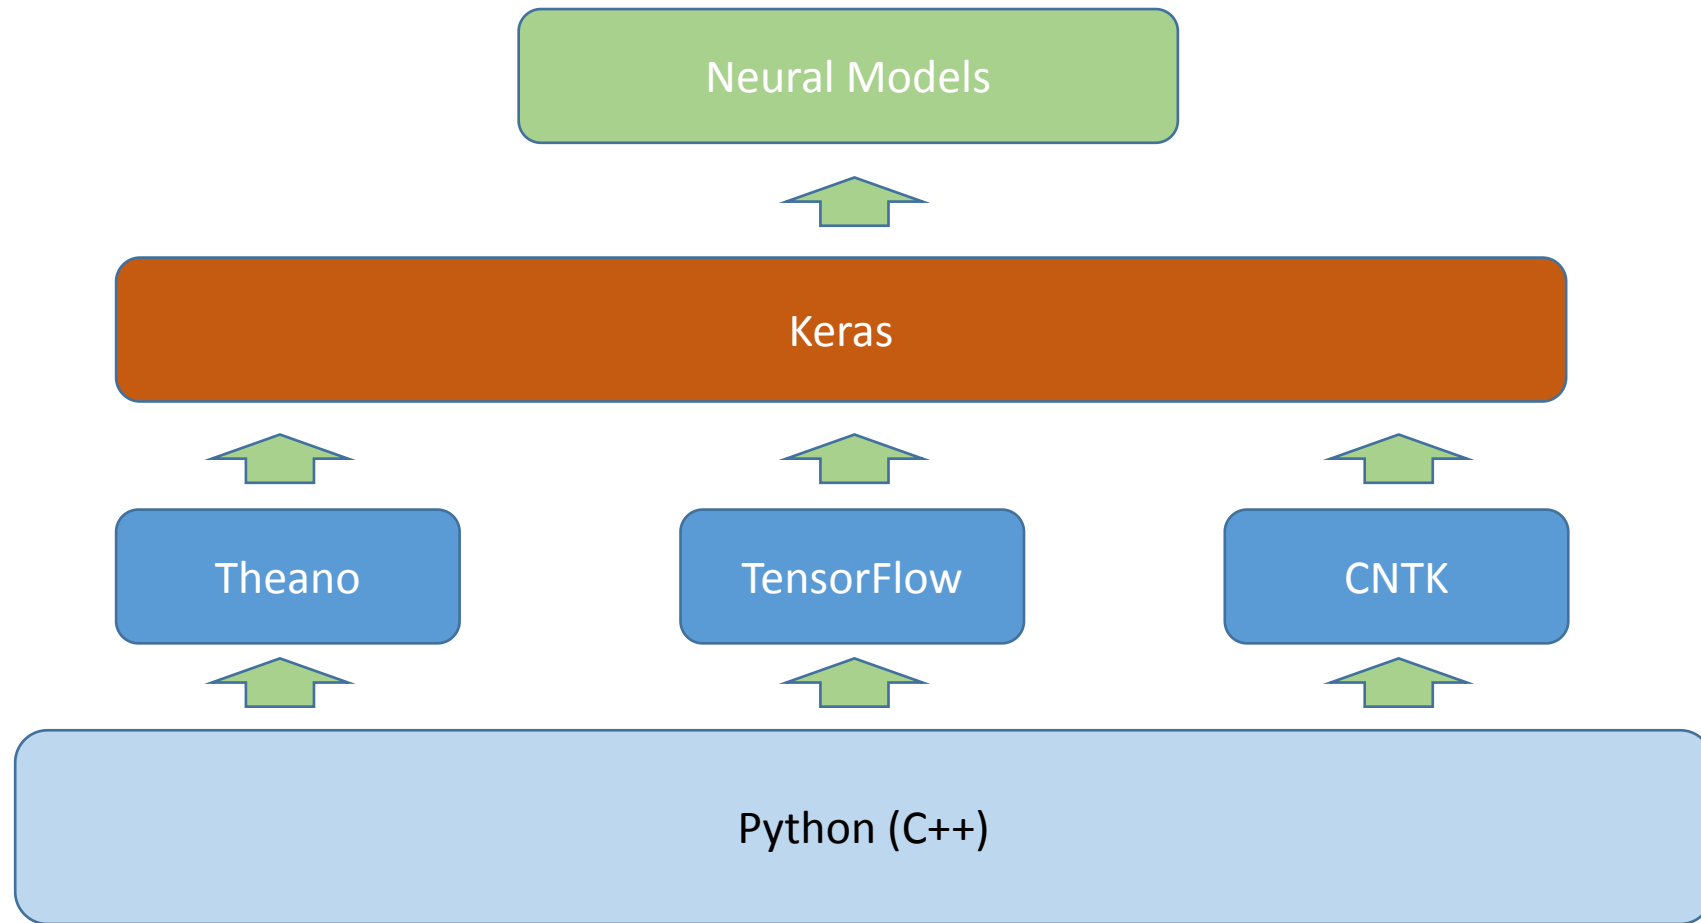
# Keras for NLP

Yan Song

# Outlines

- Keras
- Fundamentals
- Typical NLP models
- Beyond

# Keras



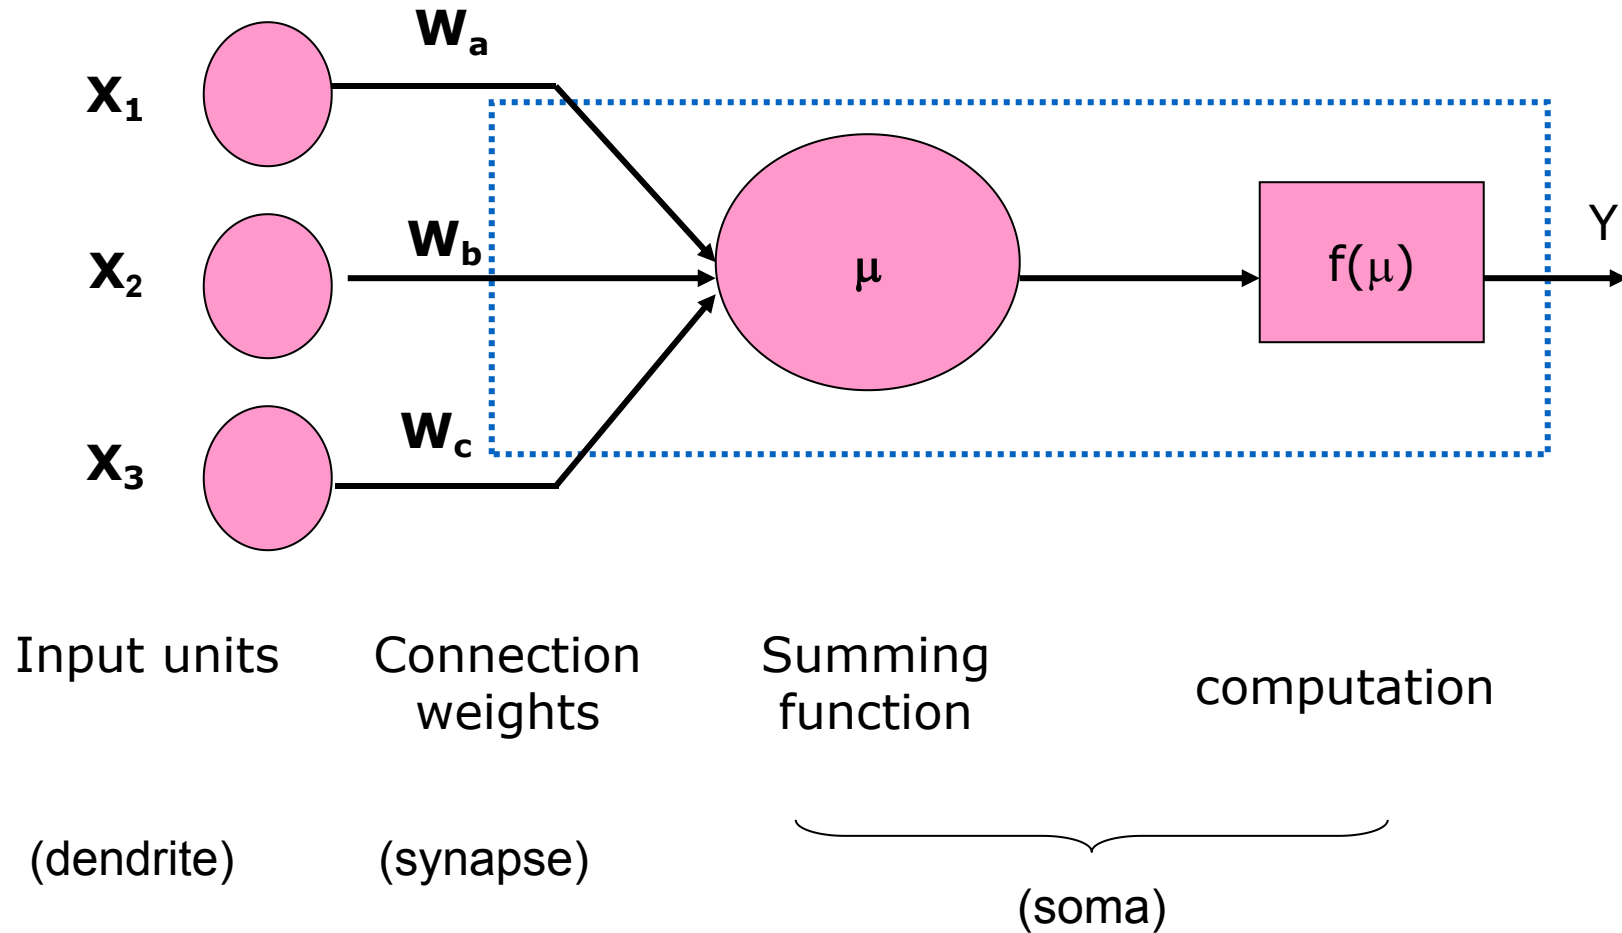
# Keras

- Features
  - EASY and FAST prototyping
  - Flow-style design
  - Seamless CPU and GPU interchanges
  - Sequential modeling or Functional API

# Fundamentals

- Symbol compilation
  - TH/TF style programming
- Tensor
  - Vector, Matrix, Tensor (3rd, 4nd, ...)
- Sequential and Functional Models

# Fundamentals




# Fundamentals

- Several important concepts
  - $W$  and  $b$
  - Layer
  - Loss function
  - Optimizer
  - Activation function

# Fundamentals

- Running Example

```
1 from keras.models import Sequential
2 from keras.layers import Dense, Activation
3
4 model = Sequential()
5 model.add(Dense(4, input_dim=3))
6 model.add(Activation('tanh'))
7 model.add(Dense(4))
8 model.add(Activation('tanh'))
9 model.add(Dense(1))
10 model.add(Activation('softmax'))
11
12 model.compile(optimizer='rmsprop',
13               loss='categorical_crossentropy',
14               metrics=['accuracy'])
15
16 model.fit(data, labels, epochs=10, batch_size=32)
```





# Fundamentals

- Core Layers
  - Dense
  - Activation
  - Dropout
  - Flatten
  - Merge
  - Embedding

# Fundamentals

- Convolutional Layers
  - Conv1D
  - Conv2D
  - Conv3D
  - MaxPooling
  - AveragePooling

# Fundamentals

- Recurrent Layers
  - SimpleRNN
  - GRU
  - LSTM

# Fundamentals

- Recurrent Layers
  - SimpleRNN
  - GRU
  - LSTM

# Fundamentals


- Typical Implementation Flow

- Model = Sequential() *#Initialize model flow*
- Model.add(...) *#Add proper layer*
- Model.compile() *#Compile the model to system level*
- Model.fit() *#Train the model*

# Typical NLP Models

- Classification - MLP (DNN)


```
36 model = Sequential()
37 model.add(Dense(512, activation='relu', input_shape=(784,)))
38 model.add(Dropout(0.2))
39 model.add(Dense(512, activation='relu'))
40 model.add(Dropout(0.2))
41 model.add(Dense(num_classes, activation='softmax'))
42
43 model.summary()
44
45 model.compile(loss='categorical_crossentropy',
46               optimizer=RMSprop(),
47               metrics=['accuracy'])
48
49 history = model.fit(x_train, y_train,
50                    batch_size=batch_size,
51                    epochs=epochs,
52                    verbose=1,
53                    validation_data=(x_test, y_test))
```



# Typical NLP Models

- Classification - CNN


```
1 model = Sequential()
2
3 model.add(Embedding(max_features,
4                     embedding_dims,
5                     input_length=maxlen))
6 model.add(Dropout(0.2))
7 model.add(Conv1D(filters,
8                 kernel_size,
9                 padding='valid',
10                activation='relu',
11                strides=1))
12 model.add(MaxPooling1D())
13 model.add(Dense(hidden_dims))
14 model.add(Dropout(0.2))
15 model.add(Activation('relu'))
16 model.add(Dense(1))
17 model.add(Activation('sigmoid'))
18
19 model.compile(loss='binary_crossentropy',
20              optimizer='adam',
21              metrics=['accuracy'])
22 model.fit(x_train, y_train,
23         batch_size=batch_size,
24         epochs=epochs,
25         validation_data=(x_test, y_test))
```



# Typical NLP Models

- Classification - LSTM

```
1 model = Sequential()
2
3 model.add(Embedding(max_features, 128))
4 model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
5 model.add(Dense(1, activation='sigmoid'))
6
7 model.compile(loss='binary_crossentropy',
8               optimizer='adam',
9               metrics=['accuracy'])
10
11 model.fit(x_train, y_train,
12          batch_size=batch_size,
13          epochs=15,
14          validation_data=(x_test, y_test))
```






# Typical NLP Models

- Sequence Tagging - LSTM

```
1 model = Sequential()  
2  
3 model.add(Embedding(max_features, 128))  
4 model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2, return_sequences=True))  
5 model.add(Dense(6, activation='softmax'))  
6  
7 model.compile(loss='binary_crossentropy',  
8               optimizer='adam',  
9               metrics=['accuracy'])  
10  
11 model.fit(x_train, y_train,  
12          batch_size=batch_size,  
13          epochs=15,  
14          validation_data=(x_test, y_test))
```



# Typical Examples


- IMDB Sentiment Classification (BiLSTM)
- IMDB Sentiment Classification (CNN)
- IMDB Sentiment Classification (FastText)

# Beyond

- Implement your own Layer
  - Protocol
  - How to do it?
  - Cautions

# Example - Attention

- Attention Model
  - Sequence
  - Weights
  - Addition



# Example - Attention

- Initialization

```
1 class SimpleAttention(Layer):
2     def __init__(self,
3                 W_regularizer=None, b_regularizer=None,
4                 W_constraint=None, b_constraint=None,
5                 bias=True, **kwargs):
6
7         """
8         Keras Layer that implements an Attention mechanism for temporal data.
9         Supports Masking.
10        Follows the work of Raffel et al. [https://arxiv.org/abs/1512.08756]
11        # Input shape
12        3D tensor with shape: `(samples, steps, features)`.
13        # Output shape
14        2D tensor with shape: `(samples, features)`.
15        :param kwargs:
16
17        Just put it on top of an RNN Layer (GRU/LSTM/SimpleRNN) with return_sequences=True.
18        The dimensions are inferred based on the output shape of the RNN.
19        Example:
20            model.add(LSTM(64, return_sequences=True))
21            model.add(Attention())
22
23        """
24        self.supports_masking = True
25        self.init = initializers.get('glorot_uniform')
26
27        self.W_regularizer = regularizers.get(W_regularizer)
28        self.b_regularizer = regularizers.get(b_regularizer)
29
30        self.W_constraint = constraints.get(W_constraint)
31        self.b_constraint = constraints.get(b_constraint)
32
33        self.bias = bias
34        super(Attention, self).__init__(**kwargs)
```

# Example - Attention


- Build Layer

```
34 def build(self, input_shape):
35     assert len(input_shape) == 3
36
37     self.W = self.add_weight((input_shape[-1],),
38                             initializer=self.init,
39                             name='{}_W'.format(self.name),
40                             regularizer=self.W_regularizer,
41                             constraint=self.W_constraint)
42
43     if self.bias:
44         self.b = self.add_weight((input_shape[1],),
45                                 initializer='zero',
46                                 name='{}_b'.format(self.name),
47                                 regularizer=self.b_regularizer,
48                                 constraint=self.b_constraint)
49
50     else:
51         self.b = None
52
53     self.built = True
```

# Example - Attention

- Compute

```
53 def compute_mask(self, input, input_mask=None):
54     # do not pass the mask to the next layers
55     return None
56
57 def call(self, x, mask=None):
58     eij = dot_product(x, self.W)
59
60     if self.bias:
61         eij += self.b
62
63     eij = K.tanh(eij)
64
65     a = K.exp(eij)
66
67     # apply mask after the exp. will be re-normalized next
68     if mask is not None:
69         # Cast the mask to floatX to avoid float64 upcasting in theano
70         a *= K.cast(mask, K.floatx())
71
72     # in some cases especially in the early stages of training the sum may be almost zero
73     # and this results in NaN's. A workaround is to add a very small positive number epsilon to the sum.
74     # a /= K.cast(K.sum(a, axis=1, keepdims=True), K.floatx())
75     a /= K.cast(K.sum(a, axis=1, keepdims=True) + K.epsilon(), K.floatx())
76
77     a = K.expand_dims(a)
78     weighted_input = x * a
79     return K.sum(weighted_input, axis=1)
80
81 def compute_output_shape(self, input_shape):
82     return input_shape[0], input_shape[-1]
```



# Conlusion

- Items we coverered:
  - Keras Basics
  - Neural Network Basics
  - Typical NLP model in Keras
  - Advances in Keras