

word2vec

Yan Song

Outline

- Neural Representation Model for Words
- word2vec
 - Concept
 - Basics
 - Details with Implementation
- Evaluation
 - Word similarity measurement

Neural Representation Model for Words

- Words are the most fundamental units in language
 - Basic semantics
- Can be used in a flexible way in many tasks
- Easy to learn
- Of course there are more studies on sub-words ;-)


Neural Representation Model for Words


- What information can be leveraged?
 - Context!
- N-gram (word association) language modeling is the most efficient way for building word-word relations
- Two ways:
 - Top-down (GloVe)
 - Bottom-up (word2vec)


word2vec - Concept


- Key idea: Predict surrounding words of every word
- Faster and can easily incorporate a new sentence/document or add a word to the vocabulary
- An indirect matrix decomposition way of learning dense representations (compared with GloVe)
- Let's start from a typical process of neural language modeling


Illustration

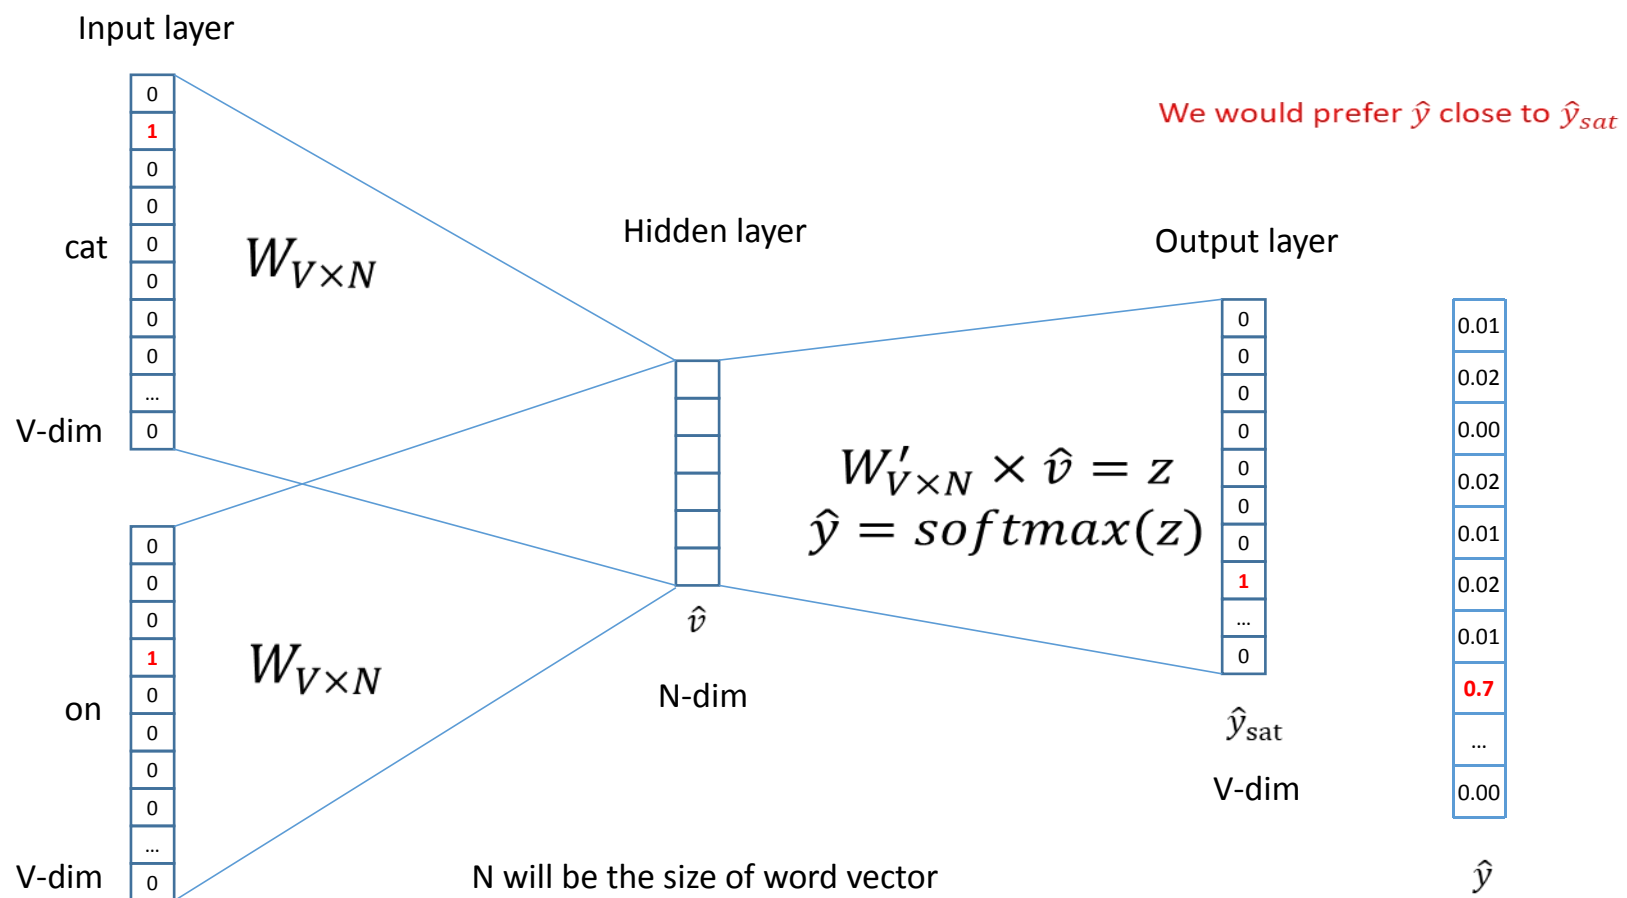











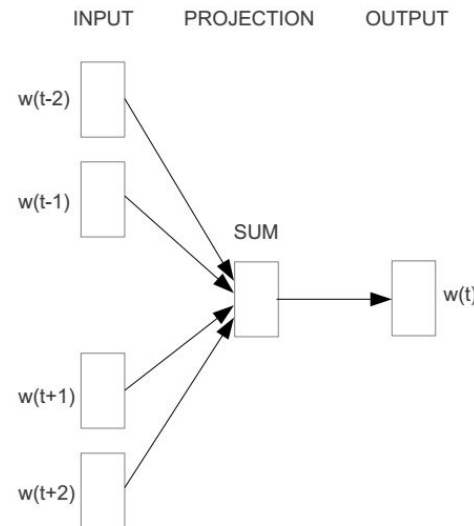





We can consider either W or W' as the word's representation. Or even take the average.

word2vec - Basics

- Two basic structures:
 - Continuous Bag of Word (CBOW): use a window of word to predict the middle word
 - Skip-gram (SG): use a word to predict the surrounding ones in window.



CBOW



Skip-gram

word2vec - Basics


- Challenges of Implementation
 - Sparse output layer
 - Matrix (vector) computation over a huge corpus
 - Deal with words with different frequencies
 - Efficient gradient updating
 - ...

word2vec - Basics

- Two prediction strategies
 - Hierarchical Softmax
 - Several pipelined softmaxes go through a tree structure
 - Effective to capture words' information when they have similar patterns in frequencies
 - Negative sampling
 - A binary classification to distinguish if a word is in context or not
 - Effective when training data is huge (less softmax operations)

Hierarchical Softmax

- An illustration
- Softmax is done at each node along the path from root to leaves
- Sibling nodes may share more frequency patterns
- Effectively convert a high-dimensional classification problem into a series of small ones



word2vec - Implementation

- Workflow
 - Load data
 - Prepare data structure
 - Training (CBOW or SG)
 - if with HS
 - predict go through hierarchical path
 - update
 - else: with NS
 - randomly choose negative samples
 - predict and update
 - Save model

word2vec - Implementation

- Load data
 - The goal is to create a vocabulary
 - Related Functions
 - AddWordToVocab
 - GetWordHash
 - LearnVocabFromTrainFile
 - ReadVocab
 - ReadWord
 - ReadWordIndex
 - ReduceVocab
 - SearchVocab
 - SortVocab

word2vec - Implementation

- Some tricks
 - All embeddings are stored in a very long 1-D array.

```
341 | a = posix_memalign((void **)&syn0, 128, (long long)vocab_size * layer1_size * sizeof(real));  
342 | if (syn0 == NULL) {printf("Memory allocation failed\n"); exit(1);}
```

- Sigmoid function is precomputed

```
695 | expTable = (real *)malloc((EXP_TABLE_SIZE + 1) * sizeof(real));  
696 | for (i = 0; i < EXP_TABLE_SIZE; i++) {  
697 |     expTable[i] = exp((i / (real)EXP_TABLE_SIZE * 2 - 1) * MAX_EXP); // Precompute the exp() table  
698 |     expTable[i] = expTable[i] / (expTable[i] + 1); // Precompute f(x) = x / (x + 1)  
699 | }
```

- important variables: **syn0**, syn1, neu1, neu1e

word2vec - Implementation

- Prepare data structure

```
195 // Create binary Huffman tree using the word counts
196 // Frequent words will have short unique binary codes
197 void CreateBinaryTree() {
198     long long a, b, i, min1i, min2i, pos1, pos2, point[MAX_CODE_LENGTH];
199     char code[MAX_CODE_LENGTH];
200     long long *count = (long long *)calloc(vocab_size * 2 + 1, sizeof(long long));
201     long long *binary = (long long *)calloc(vocab_size * 2 + 1, sizeof(long long));
202     long long *parent_node = (long long *)calloc(vocab_size * 2 + 1, sizeof(long long));
203     for (a = 0; a < vocab_size; a++) count[a] = vocab[a].cn;
204     for (a = vocab_size; a < vocab_size * 2; a++) count[a] = 1e15;
205     pos1 = vocab_size - 1;
206     pos2 = vocab_size;
207     // Following algorithm constructs the Huffman tree by adding one node at a time
208     for (a = 0; a < vocab_size - 1; a++) {
209         // First, find two smallest nodes 'min1, min2'
210         if (pos1 >= 0) {
211             if (count[pos1] < count[pos2]) {
212                 min1i = pos1;
213                 pos1--;
214             } else {
215                 min1i = pos2;
216                 pos2++;
217             }
218         } else {
219             min1i = pos2;
220             pos2++;
221         }
222         if (pos1 >= 0) {
223             if (count[pos1] < count[pos2]) {
224                 min2i = pos1;
225                 pos1--;
226             } else {
227                 min2i = pos2;
228                 pos2++;
```

```
229     }
230     } else {
231         min2i = pos2;
232         pos2++;
233     }
234     count[vocab_size + a] = count[min1i] + count[min2i];
235     parent_node[min1i] = vocab_size + a;
236     parent_node[min2i] = vocab_size + a;
237     binary[min2i] = 1;
238 }
239 // Now assign binary code to each vocabulary word
240 for (a = 0; a < vocab_size; a++) {
241     b = a;
242     i = 0;
243     while (1) {
244         code[i] = binary[b];
245         point[i] = b;
246         i++;
247         b = parent_node[b];
248         if (b == vocab_size * 2 - 2) break;
249     }
250     vocab[a].codelen = i;
251     vocab[a].point[0] = vocab_size - 2;
252     for (b = 0; b < i; b++) {
253         vocab[a].code[i - b - 1] = code[b];
254         vocab[a].point[i - b] = point[b] - vocab_size;
255     }
256 }
257 free(count);
258 free(binary);
259 free(parent_node);
260 }
261 }
```

word2vec - Implementation

- CBOW

```
422 if (cbow) { //train the cbow architecture
423     // in -> hidden
424     cw = 0;
425     for (a = b; a < window * 2 + 1 - b; a++) if (a != window) {
426         c = sentence_position - window + a;
427         if (c < 0) continue;
428         if (c >= sentence_length) continue;
429         last_word = sen[c];
430         if (last_word == -1) continue;
431         for (c = 0; c < layer1_size; c++) neul[c] += syn0[c + last_word * layer1_size];
432         cw++;
433     }
434     if (cw) {
435         for (c = 0; c < layer1_size; c++) neul[c] /= cw;
436         if (hs) for (d = 0; d < vocab[word].codelen; d++) {
437             f = 0;
438             l2 = vocab[word].point[d] * layer1_size;
439             // Propagate hidden -> output
440             for (c = 0; c < layer1_size; c++) f += neul[c] * syn1[c + l2];
441             if (f <= -MAX_EXP) continue;
442             else if (f >= MAX_EXP) continue;
443             else f = expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))];
444             // 'g' is the gradient multiplied by the learning rate
445             g = (1 - vocab[word].code[d] - f) * alpha;
446             // Propagate errors output -> hidden
447             for (c = 0; c < layer1_size; c++) neule[c] += g * syn1[c + l2];
448             // Learn weights hidden -> output
449             for (c = 0; c < layer1_size; c++) syn1[c + l2] += g * neul[c];
450         }
451         // NEGATIVE SAMPLING
452         if (negative > 0) for (d = 0; d < negative + 1; d++) {
453             if (d == 0) {
454                 target = word;
```

Projection

```
455     label = 1;
456 } else {
457     next_random = next_random * (unsigned long long)25214903917 + 11;
458     target = table[(next_random >> 16) % table_size];
459     if (target == 0) target = next_random % (vocab_size - 1) + 1;
460     if (target == word) continue;
461     label = 0;
462 }
463 l2 = target * layer1_size;
464 f = 0;
465 for (c = 0; c < layer1_size; c++) f += neul[c] * synlneg[c + l2];
466 if (f > MAX_EXP) g = (label - 1) * alpha;
467 else if (f < -MAX_EXP) g = (label - 0) * alpha;
468 else g = (label - expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))]) * alpha;
469 for (c = 0; c < layer1_size; c++) neule[c] += g * synlneg[c + l2];
470 for (c = 0; c < layer1_size; c++) synlneg[c + l2] += g * neul[c];
471 }
472 // hidden -> in
473 for (a = b; a < window * 2 + 1 - b; a++) if (a != window) {
474     c = sentence_position - window + a;
475     if (c < 0) continue;
476     if (c >= sentence_length) continue;
477     last_word = sen[c];
478     if (last_word == -1) continue;
479     for (c = 0; c < layer1_size; c++) syn0[c + last_word * layer1_size] += neule[c];
480 }
481 }
```


word2vec - Implementation

- SG

```
482 } else { //train skip-gram
483     for (a = b; a < window * 2 + 1 - b; a++) if (a != window) {
484         c = sentence_position - window + a;
485         if (c < 0) continue;
486         if (c >= sentence_length) continue;
487         last_word = sen[c];
488         if (last_word == -1) continue;
489         l1 = last_word * layer1_size;
490         for (c = 0; c < layer1_size; c++) neule[c] = 0;
491         // HIERARCHICAL SOFTMAX
492         if (hs) for (d = 0; d < vocab[word].codelen; d++) {
493             f = 0;
494             l2 = vocab[word].point[d] * layer1_size;
495             // Propagate hidden -> output
496             for (c = 0; c < layer1_size; c++) f += syn0[c + l1] * syn1[c + l2];
497             if (f <= -MAX_EXP) continue;
498             else if (f >= MAX_EXP) continue;
499             else f = expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))];
500             // 'g' is the gradient multiplied by the learning rate
501             g = (1 - vocab[word].code[d] - f) * alpha;
502             // Propagate errors output -> hidden
503             for (c = 0; c < layer1_size; c++) neule[c] += g * syn1[c + l2];
504             // Learn weights hidden -> output
505             for (c = 0; c < layer1_size; c++) syn1[c + l2] += g * syn0[c + l1];
506         }
507         // NEGATIVE SAMPLING
508         if (negative > 0) for (d = 0; d < negative + 1; d++) {
509             if (d == 0) {
510                 target = word;
511                 label = 1;
512             } else {
513                 next_random = next_random * (unsigned long long)25214903917 + 11;
514                 target = table[(next_random >> 16) % table_size];
```

Projection

```
515         if (target == 0) target = next_random % (vocab_size - 1) + 1;
516         if (target == word) continue;
517         label = 0;
518     }
519     l2 = target * layer1_size;
520     f = 0;
521     for (c = 0; c < layer1_size; c++) f += syn0[c + l1] * syn1neg[c + l2];
522     if (f > MAX_EXP) g = (label - 1) * alpha;
523     else if (f < -MAX_EXP) g = (label - 0) * alpha;
524     else g = (label - expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))]) * alpha;
525     for (c = 0; c < layer1_size; c++) neule[c] += g * syn1neg[c + l2];
526     for (c = 0; c < layer1_size; c++) syn1neg[c + l2] += g * syn0[c + l1];
527 }
528 // Learn weights input -> hidden
529 for (c = 0; c < layer1_size; c++) syn0[c + l1] += neule[c];
530 }
531 }
532 sentence_position++;
533 if (sentence_position >= sentence_length) {
534     sentence_length = 0;
535     continue;
536 }
537 }
```

word2vec - Implementation

- A closer look of HS

```
491 // HIERARCHICAL SOFTMAX
492 if (hs) for (d = 0; d < vocab[word].codelen; d++) {
493     f = 0;
494     l2 = vocab[word].point[d] * layer1_size;
495     // Propagate hidden -> output
496     for (c = 0; c < layer1_size; c++) f += syn0[c + l1] * syn1[c + l2];
497     if (f <= -MAX_EXP) continue;
498     else if (f >= MAX_EXP) continue;
499     else f = expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))];
500     // 'g' is the gradient multiplied by the learning rate
501     g = (1 - vocab[word].code[d] - f) * alpha;
502     // Propagate errors output -> hidden
503     for (c = 0; c < layer1_size; c++) neule[c] += g * syn1[c + l2];
504     // Learn weights hidden -> output
505     for (c = 0; c < layer1_size; c++) syn1[c + l2] += g * syn0[c + l1];
506 }
```

Go through the path

word2vec - Implementation

- A closer look of NS

```
507 // NEGATIVE SAMPLING
508 if (negative > 0) for (d = 0; d < negative + 1; d++) {
509     if (d == 0) {
510         target = word;
511         label = 1;
512     } else {
513         next_random = next_random * (unsigned long long)25214903917 + 11;
514         target = table[(next_random >> 16) % table_size];
515         if (target == 0) target = next_random % (vocab_size - 1) + 1;
516         if (target == word) continue;
517         label = 0;
518     }
519     l2 = target * layer1_size;
520     f = 0;
521     for (c = 0; c < layer1_size; c++) f += syn0[c + l1] * syn1neg[c + l2];
522     if (f > MAX_EXP) g = (label - 1) * alpha;
523     else if (f < -MAX_EXP) g = (label - 0) * alpha;
524     else g = (label - expTable[(int)((f + MAX_EXP) * (EXP_TABLE_SIZE / MAX_EXP / 2))]) * alpha;
525     for (c = 0; c < layer1_size; c++) neule[c] += g * syn1neg[c + l2];
526     for (c = 0; c < layer1_size; c++) syn1neg[c + l2] += g * syn0[c + l1];
527 }
```

Positive

Negative

word2vec - Implementation

- Update the input embeddings
 - CBOW

```
472 // hidden -> in
473 for (a = b; a < window * 2 + 1 - b; a++) if (a != window) {
474     c = sentence_position - window + a;
475     if (c < 0) continue;
476     if (c >= sentence_length) continue;
477     last_word = sen[c];
478     if (last_word == -1) continue;
479     for (c = 0; c < layer1_size; c++) syn0[c + last_word * layer1_size] += neule[c];
480 }
```

- SG

```
528 // Learn weights input -> hidden
529 for (c = 0; c < layer1_size; c++) syn0[c + 11] += neule[c];
```

Evaluation

- We already know cosine similarity is the main metric for intrinsic evaluation
- But, how to systematically evaluate?
 - Prepare a human annotated similarity/relatedness dataset
 - A list of word pairs
 - Compute the similarities of all pairs using the resulted embeddings
 - Compute the correlation between human annotations and the similarities

Evaluation

- Spearman's correlation

$$\rho = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2 \sum_i (y_i - \bar{y})^2}}.$$

tiger cat	7.35	0.517977544
tiger tiger	10.00	1.0
plane car	5.77	0.217732013636
train car	6.31	0.197087634791
television radio	6.77	0.257970738766
media radio	7.42	0.212014745796
bread butter	6.19	0.7017263618
cucumber potato	5.92	0.396596853365
doctor nurse	7.00	0.382799131727
professor doctor	6.62	0.288006966845
student professor	6.81	0.19973114775
smart stupid	5.81	0.228546918905
wood forest	7.73	0.229380996826

x

y

Hw4

- Train your word2vec embedding using the Text8 corpus
 - with HS (CBOW, SG)
 - with NS (CBOW, SG)
- Evaluate using spearman's correlation
- Submit a result.txt to me with
 - The results from five embeddings (4 word2vec and 1 GloVe)
 - Your observations on HS v.s. NS on CBOW and SG