



# ELMo

AI2, 2018

Lu Liu, Peter Schoener, Einar Horn



# Motivation

---



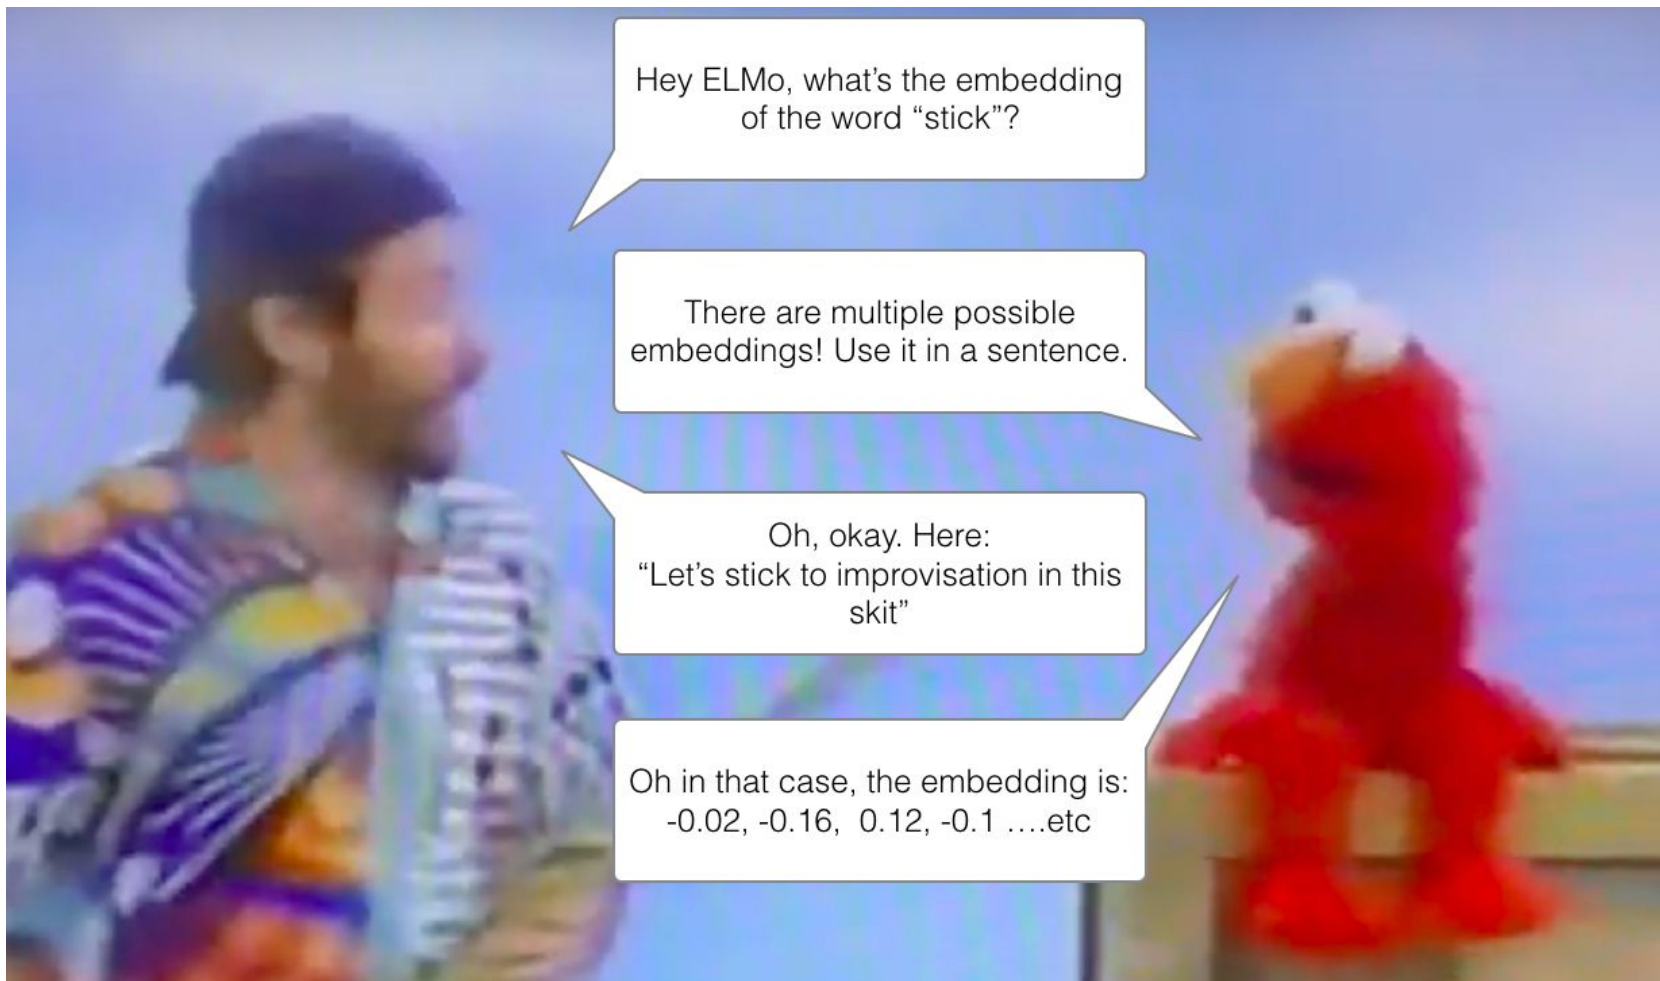
# ELMo

- ELMo - Deep Contextualised Word Representations
  - Embeddings from Language Models
- 2018, AllenNLP
- Deep bi-directional LSTM model to create word representations
- Uses a word **and** context to create an embedding



# Why not use word embeddings?

- Word embeddings (word2vec/GloVe) are like a dictionary
  - maps string -> vector
- But words can have multiple different meanings
  - **suit** yourself
  - wore a **suit**
- A word's meaning changes on **context**
- Solution: contextualized word-embeddings



Hey ELMo, what's the embedding of the word "stick"?

There are multiple possible embeddings! Use it in a sentence.

Oh, okay. Here:  
"Let's stick to improvisation in this skit"

Oh in that case, the embedding is:  
-0.02, -0.16, 0.12, -0.1 ....etc



## Pretrained Model

- Contains a language model trained on the 1B word benchmark (800 million unique tokens)
  - 93.6 million parameters
- Then we fix LM weights, fine-tune additional parameters for a specific NLP task



# Why pretraining?

- Why separate into two training phases?
- Pretraining is done on a large **unlabelled** corpus
  - General task of language modelling
    - I.e. what's the next word given the previous words?
  - Model gains a fundamental understanding of the language (syntax/semantics)
- Fine-tuning is done on a smaller labelled corpus
  - Teach the model how it should apply it's understanding of language
    - Should it be focussing on sentiment? Sentence similarity? Etc.

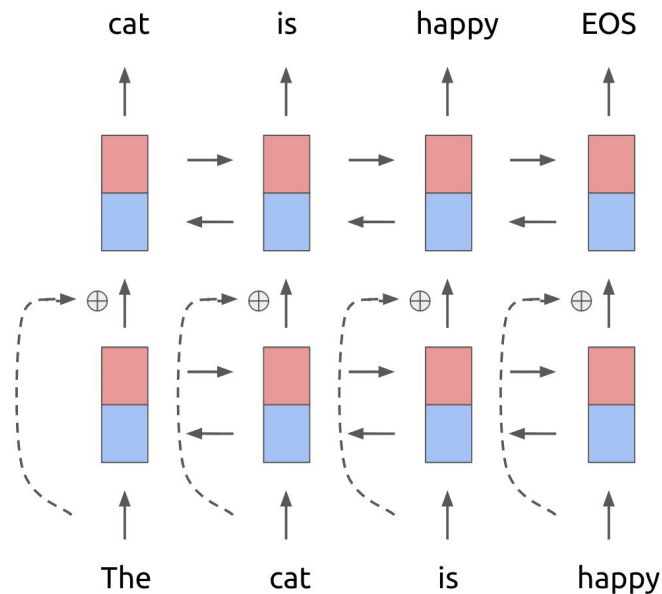
# Model-Design

---



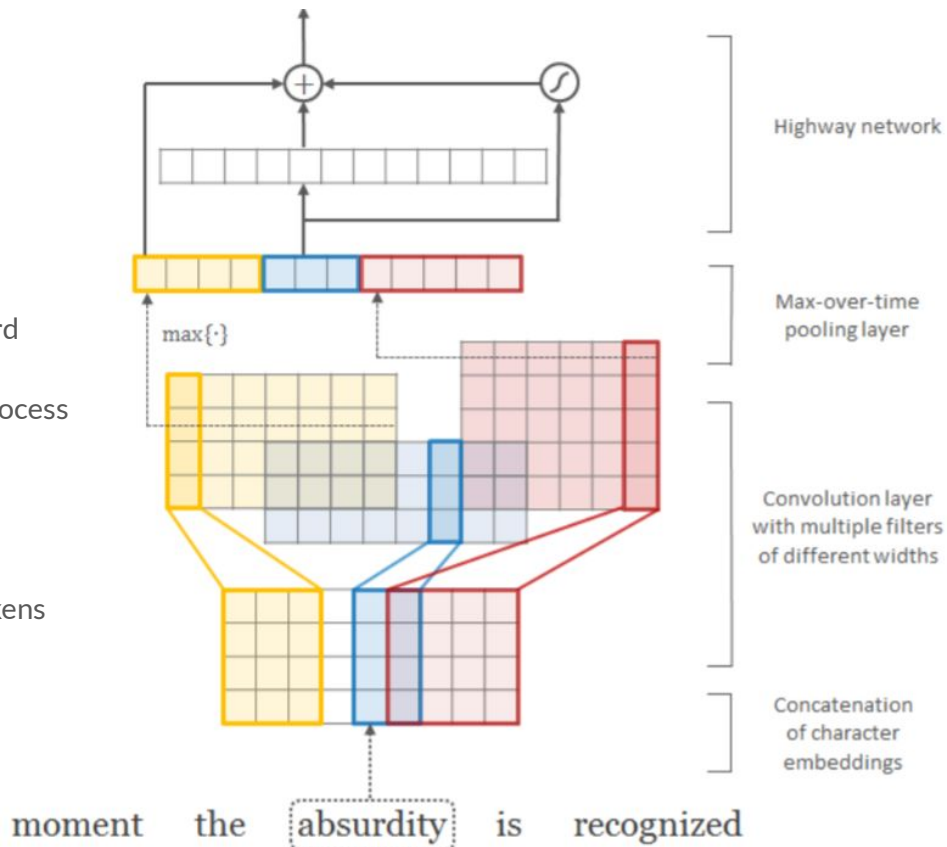
# ELMo Architecture

- ELMo architecture contains a 2 layer bi-LSTM language model
  - language model computes the probability of a word, given some prior history + future words
  - Trained on large unsupervised corpus
  - Higher layer learns semantics, lower layer learns syntax

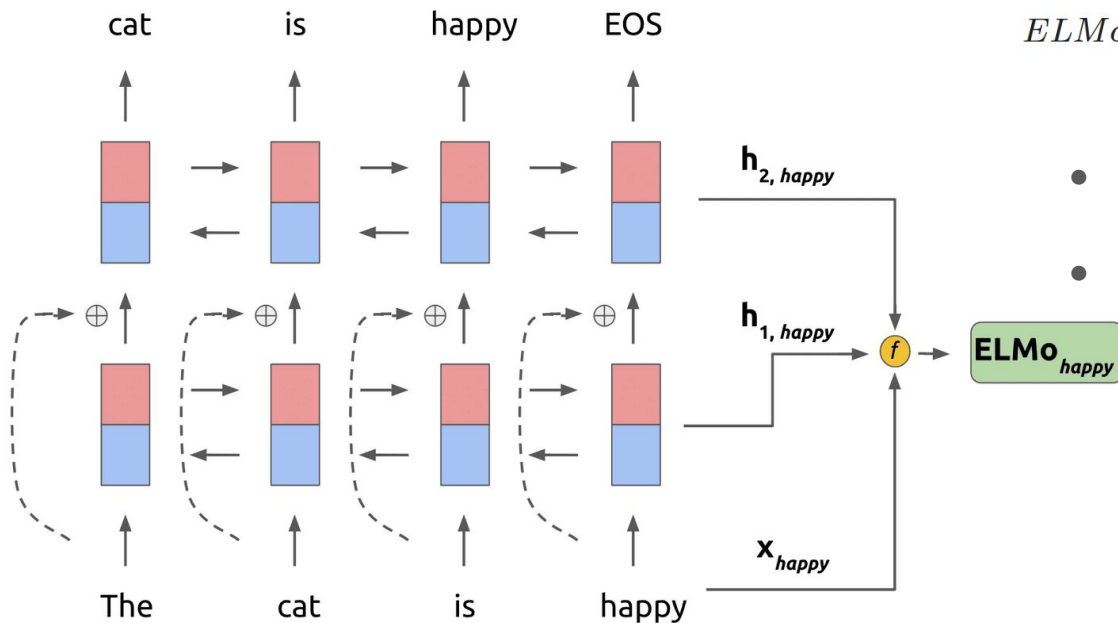


# ELMo Input

- Input layer
  - Until now, we've used 1H embeddings or word embeddings
  - ELMo uses a special **character** embedding process
  - Character embeddings
    - Capture morphological information
      - i.e. subword information
      - e.g. **pre-determine**
    - Still creates embeddings for OOV tokens



# ELMo Architecture



$$ELMo_k^{task} = \gamma_k \cdot (s_0^{task} \cdot x_k + s_1^{task} \cdot h_{1,k} + s_2^{task} \cdot h_{2,k})$$

- Task specific weights are learned by training on specific task
- Word embedding is a linear combination of each layer's output

# Model-How to learn it

---



# biLM

- Forward LM

$$p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k \mid t_1, t_2, \dots, t_{k-1}).$$

- A context-independent token representation  $\mathbf{x}_k^{LM}$ , then pass it through L layers of forward LSTMs.
- At each position k, each LSTM layer outputs a context-dependent representation

$$\vec{\mathbf{h}}_{k,j}^{LM} \text{ where } j = 1, \dots, L$$

- The top layer LSTM output  $\vec{\mathbf{h}}_{k,L}^{LM}$  is used to predict the next token  $t_{k+1}$  with a Softmax layer.



# biLM

- Backward LM
- Similar to a forward LM, except it runs over the sequence in reverse, predicting the previous token given the future context:

$$p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k \mid t_{k+1}, t_{k+2}, \dots, t_N).$$

- It can be implemented in an analogous way to a forward LM, with each backward LSTM layer  $j$  in an  $L$  layer deep model producing representations  $\overleftarrow{\mathbf{h}}_{k,j}^{LM}$  of  $t_k$  given  $(t_{k+1}, \dots, t_N)$



# Objectives

- A biLM combines both a forward and backward LM. The ELMo formulation jointly maximizes the log likelihood of the forward and backward directions:

$$\sum_{k=1}^N ( \log p(t_k | t_1, \dots, t_{k-1}; \Theta_x, \vec{\Theta}_{LSTM}, \Theta_s) \\ + \log p(t_k | t_{k+1}, \dots, t_N; \Theta_x, \overleftarrow{\Theta}_{LSTM}, \Theta_s) )$$

- The parameters of token representation  $\Theta_x$  and Softmax layer  $\Theta_s$  are tied in forward and backward directions while the parameters of LSTM are maintained separate.
- ELMo word representations are computed on top of two-layer biLMs with character convolutions, as a linear function of the internal network states.



# ELMo

- For each token  $t_k$ , an L-layer biLM computes a set of  $2L + 1$  representations

$$\begin{aligned} R_k &= \{\mathbf{x}_k^{LM}, \vec{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L\} && \text{where } \mathbf{h}_{k,0}^{LM} \text{ is the token layer and } \mathbf{h}_{k,j}^{LM} = \\ &= \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L\}, && [\vec{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM}], \text{ for each biLSTM layer.} \end{aligned}$$

- For inclusion in a downstream model, ELMo collapses all layers in  $R$  into a single vector,

$$\mathbf{ELMo}_k = E(R_k; \Theta_e)$$

- In the simplest case, ELMo just selects the top layer  $E(R_k) = \mathbf{h}_{k,L}^{LM}$  (e.g. TagLM, CoVe).





# ELMo

- More generally, ELMo computes a task specific weighting of all biLM layers:

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

- $s_j^{task}$  are softmax-normalized weights.
- $\gamma^{task}$  is the scalar parameter allows the task model to scale the entire ELMo vector.
- $\gamma$  can aid the optimization process, in some cases it also help to apply layer normalization to each biLM layer before weighting considering that the activations of each biLM layer have a different distribution,



# Using biLMs for supervised NLP tasks

- First, simply run the biLM and record all of the layer representations for each word.
- Then, we let the end task model learn a linear combination of these representations.
  - First consider the lowest layers. Given a sequence of tokens  $(t_1, \dots, t_N)$ , form a context-independent token representation  $x_k$  for each token position (pre-trained word embeddings, optionally character-based representations).
  - Then, the model forms a context-sensitive representation  $h_k$  (bidirectional RNNs, CNNs, or feed forward networks).
  - Add ELMo to the supervised model.
    - first freeze the weights of the biLM
    - then concatenate the ELMo vector  $\mathbf{ELMo}_k^{task}$  with  $x_k$
    - pass the ELMo enhanced representation  $[x_k; \mathbf{ELMo}_k^{task}]$  into the task RNN.

# Usage-Performance

---

## Model testing — results

| TASK  | PREVIOUS SOTA        |                  | OUR<br>BASELINE | ELMo +<br>BASELINE |
|-------|----------------------|------------------|-----------------|--------------------|
| SQuAD | Liu et al. (2017)    | 84.4             | 81.1            | 85.8               |
| SNLI  | Chen et al. (2017)   | 88.6             | 88.0            | $88.7 \pm 0.17$    |
| SRL   | He et al. (2017)     | 81.7             | 81.4            | 84.6               |
| Coref | Lee et al. (2017)    | 67.2             | 67.2            | 70.4               |
| NER   | Peters et al. (2017) | $91.93 \pm 0.19$ | 90.15           | $92.22 \pm 0.10$   |
| SST-5 | McCann et al. (2017) | 53.7             | 51.4            | $54.7 \pm 0.5$     |



# Model testing

- Tested on various standardized semantic tasks
  - question answering
  - entailment
  - semantic role labelling
  - coreference resolution
  - named entity extraction
  - sentiment analysis
- Direct integration in otherwise unmodified state-of-the-art methods was consistently effective
- Effects are small but consistent
- More accurate embedding modelling



# Component testing

- Semantic layer is competitive with state-of-the-art for word sense disambiguation
- Syntactic layer is better than semantic layer and CoVe at POS tagging

| Model                      | F <sub>1</sub> |
|----------------------------|----------------|
| WordNet 1st Sense Baseline | 65.9           |
| Raganato et al. (2017a)    | 69.9           |
| Iacobacci et al. (2016)    | <b>70.1</b>    |
| CoVe, First Layer          | 59.4           |
| CoVe, Second Layer         | 64.7           |
| biLM, First layer          | 67.4           |
| biLM, Second layer         | 69.0           |

Table 5: All-words fine grained WSD F<sub>1</sub>. For CoVe and the biLM, we report scores for both the first and second layer biLSTMs.

| Model                   | Acc.        |
|-------------------------|-------------|
| Collobert et al. (2011) | 97.3        |
| Ma and Hovy (2016)      | 97.6        |
| Ling et al. (2015)      | <b>97.8</b> |
| CoVe, First Layer       | 93.3        |
| CoVe, Second Layer      | 92.8        |
| biLM, First Layer       | 97.3        |
| biLM, Second Layer      | 96.8        |

Table 6: Test set POS tagging accuracies for PTB. For CoVe and the biLM, we report scores for both the first and second layer biLSTMs.

# Discussion

---



# Discussion

- Directly useful drop-in replacement
- Captures semantics very well
- Separates embeddings into meaningful components
- Fast to train even purpose-built models due to pretraining





# Reference Paper

- Peters, Matthew E., et al. "Deep contextualized word representations." arXiv preprint arXiv:1802.05365 (2018).